

Transfer_Learning_assignment

May 5, 2021

0.0.1 Model-1

0.0.2 Model-2

0.0.3 Model-3

```
[ ]: !curl --header "Host: doc-10-84-docs.googleusercontent.com" --header
↳ "User-Agent: Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36
↳ (KHTML, like Gecko) Chrome/83.0.4103.61 Safari/537.36" --header "Accept:
↳ text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/
↳ apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header
↳ "Accept-Language: en-IN,en-US;q=0.9,en;q=0.8" --header "Referer: https://
↳ drive.google.com/drive/folders/1eY4pTqtGtXSAYVzaHwWAK2Cq4bHb1Zkq" --header
↳ "Cookie: AUTH_p6auoatqfs07jahuj61ftbh1fhhqke22_nonce=91cot1ctonov8" --header
↳ "Connection: keep-alive" "https://doc-10-84-docs.googleusercontent.com/docs/
↳ securesc/hefirtae5da7uudg7909tvuvbvujep2jn/qss3i0qi2dq1k9h29322n5vgekgr8vip/
↳ 1591416900000/00484516897554883881/01214842798349466566/
↳ 1Z4TyI7FcFVEx8qdl4j09qxvxaqLSqoEu?
↳ e=download&authuser=0&nonce=91cot1ctonov8&user=01214842798349466566&hash=32f1oj749ql2aefoi
↳ -L -o 'rvl-cdip.rar'
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100 4444M	0 4444M	0 0	50.8M 0	--:--:--	0:01:27	--:--:--	42.8M

```
[ ]: get_ipython().system_raw("unrar x rvl-cdip.rar")
```

```
[ ]: !ls
```

data_final labels_final.csv rvl-cdip.rar sample_data

```
[ ]: import tensorflow as tf
import os
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import
↳ Dense, Input, Conv2D, MaxPool2D, Activation, Dropout, Flatten
```

```
from tensorflow.keras.models import Model, Sequential
import math
from tensorflow.keras.callbacks import TensorBoard
```

```
[ ]: from tensorflow.keras.utils import plot_model
```

```
[ ]: data = pd.read_csv("labels_final.csv")
data.head()
```

```
[ ]:
      path  label
0  imagesv/v/o/h/voh71d00/509132755+-2755.tif      3
1    imagesl/l/x/t/lxt19d00/502213303.tif      3
2    imagesx/x/e/d/xed05a00/2075325674.tif      2
3  imageso/o/j/b/ojb60d00/517511301+-1301.tif      3
4    imagesq/q/z/k/qzk17e00/2031320195.tif      7
```

```
[ ]: data['label'] = data['label'].astype(str)
```

```
[ ]: train, test = train_test_split(data, test_size=0.33, random_state = 42)
```

```
[ ]: print('train_size', train.shape)
      print('test_size', test.shape)
```

```
train_size (32160, 2)
test_size (15840, 2)
```

```
[ ]: tf.__version__
```

```
[ ]: '2.2.0'
```

```
[ ]: ###Image data Generator class
datagen = tf.keras.preprocessing.image.ImageDataGenerator()

##We are fitting the data to Image data generator.
ImageGenerator_train = datagen.flow_from_dataframe(dataframe = train,
    ↪, directory="data_final", x_col="path",
                                     y_col = "label",
                                     class_mode = "categorical",
    ↪
    ↪target_size=(224,224), batch_size=64)
```

Found 32160 validated image filenames belonging to 16 classes.

```
[ ]: ##We are fitting the data to Image data generator.
ImageGenerator_test = datagen.flow_from_dataframe(dataframe = test,
    ↪, directory="data_final", x_col="path",
                                     y_col = "label",
```

```

class_mode = "categorical",
↳target_size=(224,224),batch_size=64)

```

Found 15840 validated image filenames belonging to 16 classes.

```
[ ]: print(type(ImageGenerator_train))
```

```
<class 'keras_preprocessing.image.dataframe_iterator.DataFrameIterator'>
```

Model 1

```
[ ]: %load_ext tensorboard
```

```
[ ]: IMG_SHAPE = (224, 224, 3)
```

```
[ ]: VGG16_model =VGG16(include_top=False,weights="imagenet",
                        input_shape=IMG_SHAPE)
```

```
[ ]: c = 0
for layers in VGG16_model.layers:
    layers.trainable = False
    c = c+1
```

```
[ ]: VGG16_model.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080

block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080

block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0

block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160

block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808

block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808

block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0

block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808

block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

```
[ ]:
```

```
[ ]: x = VGG16_model.get_layer('block5_pool').output
x =
↳ Conv2D(filters=32,kernel_size=(3,3),strides=(1,1),padding='valid',activation='relu',kernel_
↳ keras.initializers.he_normal(seed=0),name='Conv1')(x)
x = MaxPool2D(pool_size=(2,2),strides=(2,2),padding='valid',name='pool1')(x)
x = Flatten(name='Flatten')(x)
x = Dense(units=1024,activation='relu',kernel_initializer=tf.keras.initializers.
↳ he_normal(seed=10),name='FC1')(x)
x = Dense(units=514,activation='relu',kernel_initializer=tf.keras.initializers.
↳ he_normal(seed=10),name='FC2')(x)
output_layer = Dense(units=16,activation='softmax',kernel_initializer=tf.keras.
↳ initializers.glorot_normal(seed=5),name='output_layer')(x)

model1 = Model(inputs=VGG16_model.input,outputs=output_layer )
```

```
[ ]: model1.summary()
```

```
Model: "model_1"
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

input_2 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
Conv1 (Conv2D)	(None, 5, 5, 32)	147488
pool1 (MaxPooling2D)	(None, 2, 2, 32)	0
Flatten (Flatten)	(None, 128)	0
FC1 (Dense)	(None, 1024)	132096
FC2 (Dense)	(None, 514)	526850

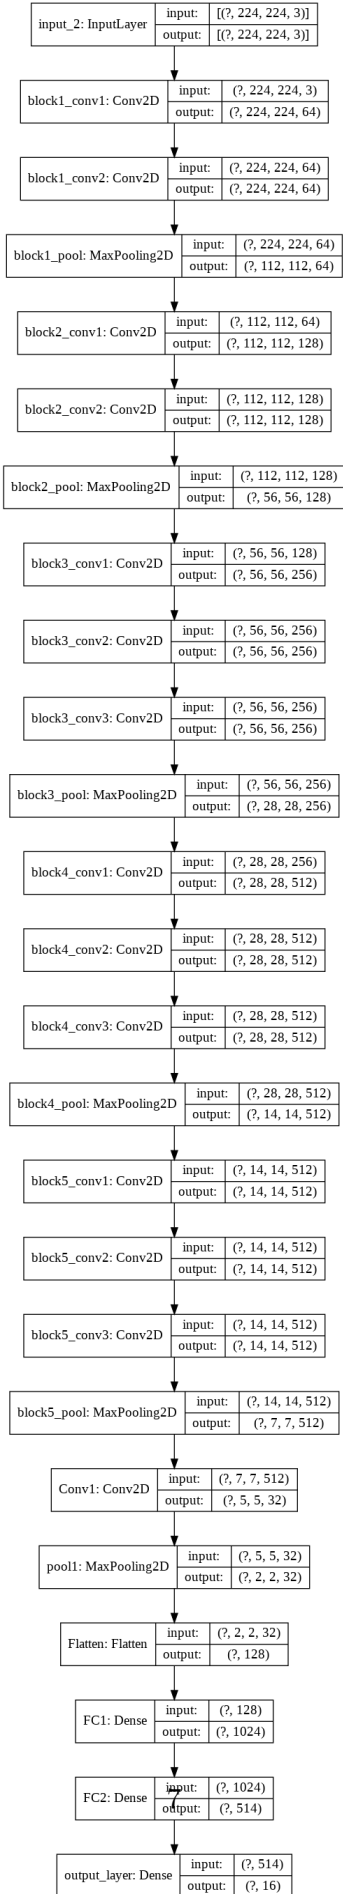
```
-----  
output_layer (Dense)          (None, 16)          8240  
=====
```

Total params: 15,529,362
Trainable params: 814,674
Non-trainable params: 14,714,688

```
-----
```

```
[ ]: plot_model(model1,to_file='model1.png',show_shapes=True,show_layer_names=True,  
              dpi=96)
```

```
[ ]:
```



```
[ ]: batch_size_points = math.ceil(32160/64)
print(batch_size_points)
```

503

```
[ ]: 15840/64
```

```
[ ]: 247.5
```

```
[ ]: log_dir = "/content/logs/"
tensorboard_callback = TensorBoard(log_dir=log_dir,histogram_freq = 1,
write_graph = True,write_grads=True)
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.

```
[ ]: model1.compile(optimizer=tf.keras.optimizers.Adam(lr=0.0001),loss='categorical_crossentropy',metrics=['accuracy'])
```

```
[ ]: model1.fit_generator(ImageGenerator_train,validation_data=ImageGenerator_test,steps_per_epoch=503,
```

Epoch 1/3

503/503 [=====] - 269s 535ms/step - loss: 1.3102 - accuracy: 0.5972 - val_loss: 1.6547 - val_accuracy: 0.5275

Epoch 2/3

503/503 [=====] - 260s 516ms/step - loss: 1.1439 - accuracy: 0.6428 - val_loss: 1.5810 - val_accuracy: 0.5539

Epoch 3/3

503/503 [=====] - 269s 535ms/step - loss: 1.0130 - accuracy: 0.6805 - val_loss: 1.5590 - val_accuracy: 0.5669

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x7f4918164668>
```

```
[ ]: %tensorboard --logdir logs
```

<IPython.core.display.Javascript object>

Observation

Using bottle feature just gives accuracy of 52% , it means dataset is not similar to imagenet or some more tuning is required

Model 2


```
[ ]: VGG16_model2 =VGG16(include_top=False,weights="imagenet",
                           input_shape=IMG_SHAPE)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58892288/58889256 [=====] - 1s 0us/step

```
[ ]: for layers in VGG16_model2.layers:
      layers.trainable= False
```

```
[ ]: VGG16_model2.summary()
```

Model: "vgg16"



Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808

block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

```
[ ]: x = VGG16_model2.get_layer('block5_pool').output

x = 
    ↳ Conv2D(4096, kernel_size=[7, 7], strides=(1, 1), padding='valid', activation='relu')(x)
x = 
    ↳ Conv2D(4096, kernel_size=[1, 1], strides=(1, 1), padding='valid', activation='relu')(x)
x = Flatten(name='flatten')(x)

output_layer = Dense(units=16, activation='softmax', kernel_initializer=tf.keras.
    ↳ initializers.glorot_normal(seed=5), name='output_layer')(x)

model2 = Model(inputs=VGG16_model2.input, outputs=output_layer )
```

```
[ ]: model2.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0

block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792

block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928

block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0

block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856

block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584

block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0

block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168

block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d (Conv2D)	(None, 1, 1, 4096)	102764544
conv2d_1 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten (Flatten)	(None, 4096)	0
output_layer (Dense)	(None, 16)	65552

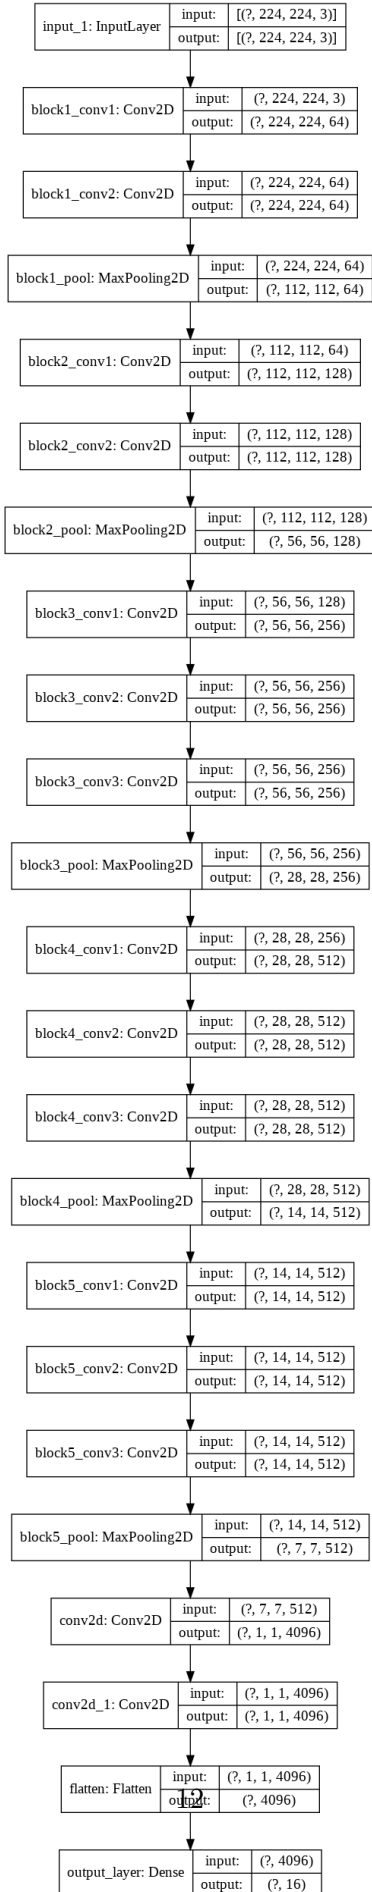
=====

Total params: 134,326,096
Trainable params: 119,611,408
Non-trainable params: 14,714,688

=====

```
[ ]: plot_model(model2,to_file='model1.png',show_shapes=True,show_layer_names=True,
              dpi=96)
```

```
[ ]:
```



```
[ ]: mkdir logs
```

```
[ ]: log_dir = "logs"
    tensorboard_callback_
    ↳TensorBoard(log_dir=log_dir,histogram_freq=1,write_graph=True,write_grads=True)
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the
`TensorBoard` Callback.

```
[ ]: model2.compile(optimizer=tf.keras.optimizers.Adam(lr=0.
    ↳001),loss='categorical_crossentropy',metrics=['accuracy'])
    model2.fit_generator(ImageGenerator_train,validation_data_
    ↳ImageGenerator_test,steps_per_epoch=503,validation_steps = 247,
    ↳epochs=3,callbacks=[tensorboard_callback])
```

WARNING:tensorflow:From <ipython-input-40-3bef14f92f4f>:2: Model.fit_generator
(from tensorflow.python.keras.engine.training) is deprecated and will be removed
in a future version.

Instructions for updating:

Please use Model.fit, which supports generators.

Epoch 1/3

503/503 [=====] - 336s 668ms/step - loss: 5.0102 -
accuracy: 0.4744 - val_loss: 1.6662 - val_accuracy: 0.5326

Epoch 2/3

503/503 [=====] - 335s 665ms/step - loss: 1.3930 -
accuracy: 0.5907 - val_loss: 1.2908 - val_accuracy: 0.5969

Epoch 3/3

503/503 [=====] - 341s 678ms/step - loss: 1.1142 -
accuracy: 0.6569 - val_loss: 1.2092 - val_accuracy: 0.6301

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x7f18900cdc18>
```

```
[ ]: !kill
```

```
[ ]: %tensorboard --logdir logs
```

<IPython.core.display.Javascript object>

```
[ ]: Here we are using covnet as
```

Observation

- 1) Here we are using FC as covnet which and training only new layers , But
still we are not very good accuracy , it means dataset is different than
imagenet dataset , and need to train some more layers .

Model 3

```
[ ]: VGG16_model3 =VGG16(include_top=False,weights="imagenet",  
                           input_shape=IMG_SHAPE)
```

```
[ ]: for layers in VGG16_model3.layers[:-6]:  
      layers.trainable = False
```

```
[ ]: VGG16_model3.summary()
```



Model: "vgg16"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808

block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 9,439,232		
Non-trainable params: 5,275,456		

```
[ ]: x = VGG16_model3.get_layer('block5_pool').output

x = 
    ↳ Conv2D(4096, kernel_size=[7, 7], strides=(1, 1), padding='valid', activation='relu')(x)
x = 
    ↳ Conv2D(4096, kernel_size=[1, 1], strides=(1, 1), padding='valid', activation='relu')(x)
x = Flatten(name='flatten')(x)

output_layer = Dense(units=16, activation='softmax', kernel_initializer=tf.keras.
    ↳ initializers.glorot_normal(seed=5), name='output_layer')(x)

model3 = Model(inputs=VGG16_model3.input, outputs=output_layer )
```

```
[ ]: model3.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #

input_2 (InputLayer)	[(None, 224, 224, 3)]	0

block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792

block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928

block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0

block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856

block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584

block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0

block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168

block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080

block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d_2 (Conv2D)	(None, 1, 1, 4096)	102764544
conv2d_3 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten (Flatten)	(None, 4096)	0
output_layer (Dense)	(None, 16)	65552

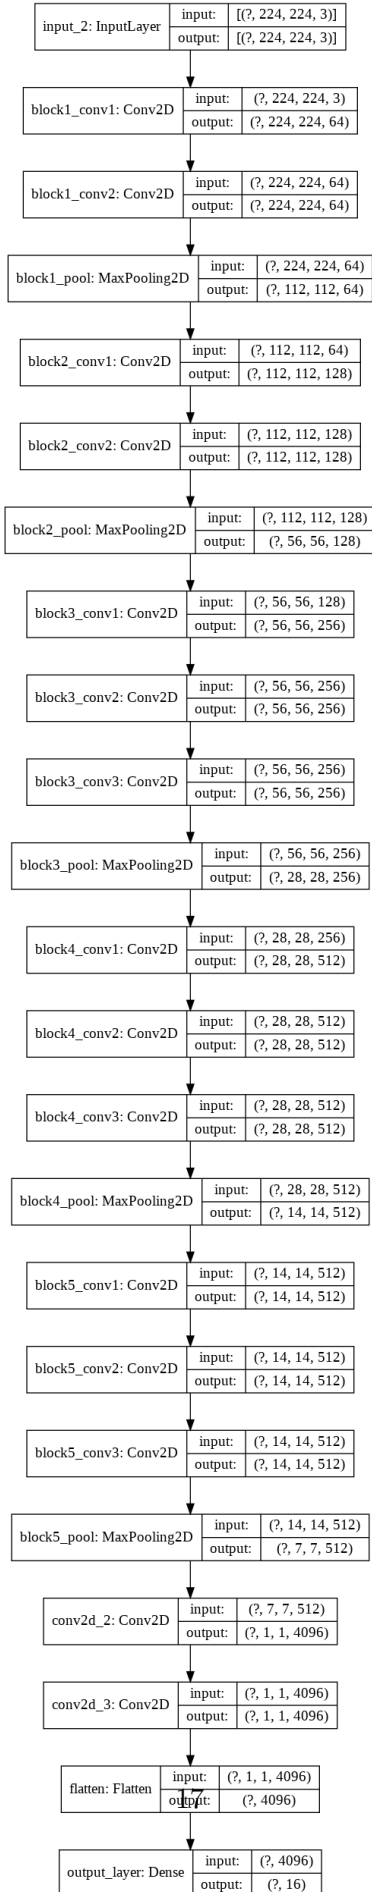
=====

Total params: 134,326,096
Trainable params: 129,050,640
Non-trainable params: 5,275,456

=====

```
[ ]: plot_model(model3,to_file='model1.png',show_shapes=True,show_layer_names=True,
            dpi=96)
```

```
[ ]:
```

```
[ ]: rm -r /content/logs/
```

```
[ ]: mkdir logs
```

```
[ ]: log_dir = "logs"
    tensorboard_callback = TensorBoard(log_dir =_
    ↪log_dir,histogram_freq=1,write_graph = True,write_grads = True)
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the
`TensorBoard` Callback.

```
[ ]: model3.compile(optimizer=tf.keras.optimizers.SGD(learning_rate=0.
    ↪0001,momentum=0.
    ↪9,nesterov=True),loss='categorical_crossentropy',metrics=['accuracy'])
    model3.fit_generator(ImageGenerator_train,validation_data =_
    ↪ImageGenerator_test,steps_per_epoch=503,validation_steps =_
    ↪247,epochs=3,callbacks = [tensorboard_callback])
```

Epoch 1/3

503/503 [=====] - 520s 1s/step - loss: 1.2265 -
accuracy: 0.6345 - val_loss: 0.9167 - val_accuracy: 0.7251

Epoch 2/3

503/503 [=====] - 527s 1s/step - loss: 0.5933 -
accuracy: 0.8200 - val_loss: 0.8845 - val_accuracy: 0.7451

Epoch 3/3

503/503 [=====] - 519s 1s/step - loss: 0.2931 -
accuracy: 0.9156 - val_loss: 0.8730 - val_accuracy: 0.7597

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x7f18426bd8d0>
```

```
[ ]: !kill
```

kill: usage: kill [-s sigspec | -n signum | -sigspec] pid | jobspec ... or kill
-l [sigspec]

```
[ ]: %tensorboard --logdir logs
```

<IPython.core.display.Javascript object>

```
[ ]:
```

Observation

- 1) Here we are making same model as model 2 but here we are training new
layers + Top most 6 layers and we are getting accuracy of 91% ,which means

data is very different than magnet .